

SEARCH · ANSWER ENGINE · CONVERSION

The Oman Authority Plan

A live audit of aiprofitlab.io and a 90-day execution roadmap to own AI-automation search in Oman — in Google, in Arabic, and inside the answers ChatGPT, Claude, Perplexity and Gemini give.

00 Contents

01	Executive summary — the one thing to fix first	The diagnosis
02	Where you actually stand today	Scorecard + evidence
03	The three structural gaps	Why traffic isn't becoming money
04	Technical audit — findings and fixes	12 findings, ranked
05	Speed and efficiency	Make it fast
06	GEO / AEO — becoming the answer LLMs give	Your stated priority
07	Ranking #1 in Google — the realistic path	EN + AR keyword plan
08	Turning visitors into buyers	Conversion architecture
09	The Proof Engine	Your binding constraint
10	Local SEO and Google Business Profile	Blocking item
11	The 90-day roadmap, week by week	Sized to 10–20 hrs/week
12	Months 4–12 and how to measure it	Beyond the sprint
A	Appendix A — copy-paste fixes	Code
B	Appendix B — keyword target list	EN + AR

HOW THIS DOCUMENT WAS MADE

Every finding here was verified against the **live site** on 25 August 2026 — HTTP headers, rendered HTML, schema payloads, asset weights, redirect behaviour and search results — not inferred from the repository. Where a claim could not be verified live it is marked as such. Numbers you can check yourself are given with the URL or file they came from.

01 Executive summary

You have built the hard part already. What is missing is not content, not technical SEO, and not effort — it is that your traffic and your selling machine are sitting in two different rooms.

aiprofitlab.io is, technically, in the top few percent of small business sites I could audit. The `robots.txt` is better than most enterprise sites — it names GPTBot, ClaudeBot, PerplexityBot, OAI-SearchBot and Google-Extended and explicitly allows each one, with a written rationale for why `noindex` is used instead of `Disallow`. The `llms.txt` is genuinely excellent and even carries brand-disambiguation language separating you from the unrelated `aiprofitlab.com` trading products. The `.htaccess` is carefully engineered, with comments explaining why each cache rule exists. You publish 154 English and 154 Arabic articles at a median of 1,713 words, with real citations to `cbo.gov.om` and `mtcit.gov.om`. Schema is deep: `Service` with real `Offer` prices, `HowTo`, `FAQPage`, `ProfessionalService`. Fonts load asynchronously and correctly. Brotli is on, assets are immutably cached, and the CDN returns edge hits at roughly 0.35s.

That work is already paying. Searching for AI automation and WhatsApp receptionist topics in Oman returns **your blog articles repeatedly** — often several in a single result set. You are winning the informational layer.

THE FINDING THAT MATTERS MOST

Your AI buyer agent, **Aiden**, is installed on 0 of your 308 blog articles. It is on the homepage and the money pages. But the blog is what ranks, and the blog is where organic visitors land. Verified live on 25 Aug 2026: `/blog/en/2026-06-16-whatsapp-ai-receptionist-cost-oman-2026/` returns 0 occurrences of `aiden-chat.js`, while `/` returns 1.

You sell a product whose entire promise is "an AI agent that answers your buyers at two in the morning." A buyer arrives from Google onto one of your articles at 11pm, reads 1,700 words that convince them, and then meets... five WhatsApp links and no agent. The single most persuasive demonstration you own — the thing working, on you, right now — is switched off on the exact pages where strangers meet you.

The five moves, in order

#	MOVE	EFFORT	WHY THIS ORDER
1	Put Aiden on all 308 articles	4–6 hrs	Biggest conversion gain available. The traffic already exists.
2	Verify the Google Business Profile	1 hr + wait	Blocks the entire local pack. Nothing local works until done.
3	301 www → apex	15 min	Your whole site is currently duplicated on www.
4	Build the Proof Engine	Ongoing	You have zero publishable proof. Everything else is capped by this.
5	Publish commercial pages, not just articles	3–4 wks	Competitors own the buying queries. Your blog owns the learning ones.

THE HONEST CONSTRAINT

You told me you have **nothing publishable** in the way of clients or testimonials. That is the ceiling on this entire plan, and no amount of SEO removes it. Section 09 is therefore not optional garnish — it is the load-bearing section of this document. A competitor, Nuqta, publishes six named clients including Oman Flour Mills. Until you can answer "who else has done this," you will lose every deal that goes to a side-by-side comparison.

02

Where you actually stand

Measured live on 25 August 2026. Green means genuinely good — not "acceptable."

332	308	1,713	0.35s	0
URLS IN SITEMAP	ARTICLES EN + AR	MEDIAN WORDS/ARTICLE	TTFB, EDGE HIT	PUBLISHED TESTIMONIALS
AREA	VERDICT	WHAT THE LIVE SITE SHOWS		
Crawler access	EXCELLENT	Explicit <code>Allow</code> for GPTBot, OAI-SearchBot, ChatGPT-User, ClaudeBot, Claude-User, Claude-SearchBot, PerplexityBot, Perplexity-User, Google-Extended, Applebot-Extended, Bingbot.		
llms.txt	EXCELLENT	Present, complete, with pricing, URL map, founder bio and brand-collision disambiguation. Very few sites in the GCC have this.		
Content depth	EXCELLENT	154 EN + 154 AR. None under 995 words. 28 over 2,000. Real gov/bank citations.		
Schema markup	STRONG	<code>Service + Offer</code> with published OMR prices, <code>HowTo</code> , <code>FAQPage</code> , <code>ProfessionalService</code> , <code>Person</code> , <code>BreadcrumbList</code> on all articles.		
hreflang	CORRECT	Self-referencing, reciprocal, with <code>x-default</code> . Verified on <code>/en/about/</code> ↔ <code>/about/</code> .		
Caching / compression	STRONG	Brotli on HTML/JS/SVG. <code>max-age=31536000</code> , <code>immutable</code> on hashed assets. CDN edge <code>HIT</code> .		
Font loading	CORRECT	preconnect + preload + <code>media="print"</code> swap + <code><noscript></code> fallback. Textbook.		
Informational ranking	WINNING	Your articles appear repeatedly for Oman AI-automation and WhatsApp queries.		
Commercial ranking	LOSING	4Ys, Inkstone, QwicLink and Nuqta surface for buying-intent service queries. Your money pages do not.		
Conversion on traffic pages	BROKEN	Aiden on 0/308 articles. No email capture. No lead magnet.		
Social proof	ABSENT	0 of 375 pages carry <code>Review</code> or <code>aggregateRating</code> . No named client anywhere.		
www duplication	OPEN	<code>www.aiprofitlab.io</code> returns HTTP 200, zero redirects . Entire site duplicated.		
Google Business Profile	UNVERIFIED	Address corrected, verification badge not obtained. Local pack unreachable until it is.		
Page weight	IMPROVABLE	1.14 MB first load. Google Analytics alone is 571 KB raw / 188 KB gzipped — 48% of it.		

READ THIS BEFORE THE CRITICAL SECTIONS

The list of problems below is long, but almost all of it is **plumbing between things you have already built** — not new construction. You are not behind. You are disconnected.

03 The three structural gaps

Gap one — the traffic and the machine are in different rooms

Your site has two populations of pages built by two different systems, and they behave differently:

PAGE TYPE	COUNT	GETS ORGANIC TRAFFIC	HAS AIDEN	CSS STRATEGY
Blog articles	308	YES	NO — 0	Shared hashed stylesheet, cacheable. Correct.
Money pages (v4)	~14	LITTLE	YES	32–66 KB inlined per page, uncacheable.

Each population has exactly what the other one needs. The articles have the traffic and the correct CSS architecture; the money pages have the conversion agent. Swapping those two things across is perhaps six hours of work and is the highest-return action available to you this quarter.

Gap two — you rank for learning, not for buying

Search results for Oman AI-automation topics return your articles again and again. Search results for *buying* — an AI automation company, a WhatsApp automation provider in Muscat — return 4Ys, Inkstone, QwicLink and Nuqta. Those competitors have thin content but strong commercial pages. You have the reverse.

You already have the raw material for the fix sitting unused in the repository: `.agents/competitors/pages/` contains drafted copy for *AI automation companies in Oman, vs hiring admin staff* and *vs QwicLink*. None of it is live. Section 07 says which to publish and in what order — and section 09 explains why the QwicLink page must wait.

Gap three — no proof, and no machine for making proof

Zero of 375 pages carry review or rating schema. There is no named client on the site. This is not a marketing oversight, it is a structural gap: there is no *process* that turns delivered work into publishable evidence. Even when you complete builds, nothing currently captures consent, before/after numbers or a quotable sentence. Section 09 builds that machine, and it must start now because proof has a lead time that no amount of budget compresses.

A NOTE ON THE PRICE OBJECTION

Your own competitor research is blunt about this and it is worth restating: QwicLink publishes OMR 15/29/49 per month, putting their AI tier near OMR 617 in year one against roughly OMR 1,400 for your Starter. **You are about 2.3× more expensive locally.** Never argue on price against local rivals — you lose that argument on the facts. Argue on done-for-you scope, on the operator-not-agency story, and on the fact that nothing stops working if the optional Growth Desk is cancelled. That last point is a genuinely strong differentiator against subscription competitors and is currently under-used.

04 Technical audit — findings and fixes

Twelve findings, ranked by revenue impact rather than by technical severity. Every one was verified against the live site.

#	FINDING	PRIORITY	EVIDENCE AND FIX
1	Aiden absent from all articles	CRITICAL	0 of 308 blog pages contain <code>aiden-chat.js</code> ; the homepage contains it. Verified live. <i>Fix</i> : extend <code>tools/add_aiden_widget.py</code> to the <code>blog/en/</code> and <code>blog/ar/</code> trees, or add the tag to the article template and re-run <code>reskin_articles.py</code> . Respect the idempotency invariants so a re-run does not stack duplicate tags.
2	No email capture anywhere	CRITICAL	No newsletter, lead magnet, or content upgrade on any article. Every visitor who is not ready to message you on WhatsApp today is lost permanently. <i>Fix</i> : one lead magnet, offered inline on articles — see section 08.
3	Zero review / rating schema	CRITICAL	<code>aggregateRating</code> : 0 files. <code>Review</code> : 0 files. Out of 375. <i>Fix</i> : section 09 — and do not fabricate these. Invented reviews are a manual-action risk and, more importantly, a lie a buyer can catch.
4	www serves a full duplicate site	HIGH	<code>https://www.aiprofitlab.io/</code> → HTTP 200, 0 redirects. All 332 URLs exist twice. Canonicals correctly point at the apex, which limits the damage, but Google still has to crawl and reconcile a second copy. <i>Fix</i> : 301 at the host/CDN layer — see Appendix A.
5	GBP not verified	HIGH	Address corrected to South Al Khuwair, Bousher; verification badge not obtained. Without it you cannot appear in the Muscat local pack or Google Maps for "AI automation near me". <i>Fix</i> : request postcard/phone verification this week — it has a multi-week lead time.
6	No breadcrumbs on money pages	HIGH	310 files carry <code>BreadcrumbList</code> — every one is a blog page or hub. The homepage, services, process, about, contact and checkout pages carry none. These are the pages whose SERP appearance matters most commercially. <i>Fix</i> : Appendix A.
7	Analytics is 48% of page weight	HIGH	<code>gtag.js</code> = 571,511 bytes raw / 187,680 gzipped, against a 1.14 MB total first load. It is by far your largest asset — larger than every image on the page combined. <i>Fix</i> : section 05.
8	Inline CSS on money pages	MEDIUM	32–66 KB of CSS inlined per money page (homepage 51,847 B = 44% of the document; Arabic homepage 66,455 B). Across the site that is 1.37 MB of CSS that no browser can ever cache. Your blog pages already solve this correctly with <code>/assets/css/apl-article.eda58827d4.css</code> . <i>Fix</i> : apply the same hashed-stylesheet pattern in <code>tools/build_v4.py</code> .
9	Image attributes incomplete	MEDIUM	Homepage: 8 <code></code> tags — 5 without <code>loading</code> , 2 without <code>width / height</code> , 2 without <code>alt</code> . Missing dimensions cause layout shift (CLS); missing alt costs accessibility and image search. <i>Fix</i> : Appendix A.
10	Inconsistent URL architecture	MEDIUM	<code>/</code> is English but <code>/about/</code> , <code>/services/</code> , <code>/process/</code> , <code>/contact/</code> are Arabic , with English at <code>/en/...</code> . So the root is English at the homepage and Arabic one level down. hreflang is correct so this is not breaking anything today, but it is confusing for humans, for link-builders, and for anyone reasoning about your site structure — including LLMs. <i>Fix</i> : leave it unless you are already planning a restructure; if you do restructure, move Arabic to <code>/ar/...</code> throughout and 301 the old paths. Do not do this during the 90-day sprint.

11	Schema scope too narrow	MEDIUM	<code>areaServed</code> lists Oman only, though you sell across the GCC. No <code>knowsAbout</code> on the organisation, which is one of the clearest topical-authority signals available to entity-based systems. <i>Fix</i> : Appendix A.
12	WAF throttles burst crawling	WATCH	During this audit, rapid sequential requests triggered Hostinger <code>403</code> responses; spaced requests returned <code>200</code> consistently. Normal browsing is unaffected, but a crawler working through 332 URLs quickly may be throttled. <i>Fix</i> : nothing to change now — but if Search Console reports crawl anomalies, this is the first place to look.

THINGS I CHECKED THAT ARE FINE — DO NOT SPEND TIME HERE

`http` → `https` redirects correctly (301). `/en/` → `/` redirects correctly (301) — a stale search-index entry for `/en/` still exists but will resolve itself. hreflang pairs are reciprocal and self-referencing with `x-default`. Fonts are loaded with the correct async pattern including a `<noscript>` fallback. Sitemap is valid, complete and dated. Canonicals are present and correct on every page checked. Brotli and immutable caching are configured properly.

05 Speed and efficiency

The site is already fast. The remaining wins are concentrated in two places, and one of them is worth more than everything else combined.

Where the 1.14 MB actually goes

ASSET	RAW	TRANSFERRED	SHARE	ACTION
gtag.js (Google Analytics)	571,511	187,680	48%	Defer or replace
Homepage HTML	116,745	~18,000 (br)	10%	Externalise CSS
aiden-chat.js	59,284	brotli	5%	Lazy-init on idle
demo-dashboard-960.webp	49,138	49,138	4%	Lazy-load
demo-whatsapp-960.webp	33,574	33,574	3%	Lazy-load
Hero posters (2 × webp)	51,746	51,746	4%	Correct — preloaded
Wordmark SVGs (2)	25,288	brotli	2%	Deduplicate

Move one — deal with Google Analytics

188 KB transferred to record page views is a poor trade, and it is the largest single cost on your site. Three options, in order of preference:

- **Replace it.** A privacy-first analytics script (Plausible, Fathom, Umami) is roughly 1–3 KB — a **98% reduction** on your heaviest asset. You lose Google Ads remarketing audiences; if you are not running Google Ads, you lose nothing you use.
- **Defer it.** Keep GA4 but load it after first interaction or on `requestIdleCallback`. Costs you a small number of bounced sessions in the data; costs the user nothing during paint.
- **Server-side tag it.** Most robust, most work. Only worth it if you commit to Google Ads later.

Given you are one operator optimising for organic and LLM visibility rather than paid media, **replace it**. Keep GA4 running in parallel for one month if you want continuity in the numbers, then drop it.

Move two — externalise the money-page CSS

Your blog already does this correctly and your money pages do not. The blog pattern — one content-hashed stylesheet, immutably cached, shared across every page — means a reader who opens three articles downloads the CSS once. On the money pages, a visitor going homepage → services → checkout downloads 51 KB, then 38 KB, then 39 KB of largely identical CSS. Applying the blog's own pattern in `tools/build_v4.py` fixes it.

Keep a small critical-CSS block inline — roughly the first 8–14 KB covering the header and hero — and move the rest out. That preserves fast first paint while making every subsequent page nearly free.

Move three — images and the Aiden widget

- Add `loading="lazy"` and `decoding="async"` to the five below-the-fold homepage images. Leave the preloaded hero posters exactly as they are — those are correctly configured and must not be lazy-loaded.
- Add explicit `width` and `height` to the two images missing them. This removes layout shift at no cost.

- Add real `alt` text to the two missing it — needed for accessibility, image search, and for LLMs parsing the page.
- When you put Aiden on 308 articles, initialise it on `requestIdleCallback` or first scroll rather than on load, so it never competes with article text for the first paint.

REALISTIC EXPECTATION

These changes will not transform your Core Web Vitals, because your vitals are already decent — TTFB around 0.35s on an edge hit, correct font loading, preloaded hero images. Expect a meaningful improvement in total bytes and in repeat-visit speed, not a dramatic score jump. **Do not spend the sprint chasing a perfect Lighthouse score.** The conversion and proof work in sections 08 and 09 is worth an order of magnitude more to revenue than the last few performance points.

06 GEO / AEO — becoming the answer

You chose LLM citation over classic ranking as a priority. That is the right call for your market, and you are closer to it than you probably think — but the last mile is not about content.

What you have already won

Most sites fail LLM citation at the first hurdle: the crawler is blocked, or there is nothing machine-readable to consume. You have cleared both. Your `robots.txt` explicitly admits every major AI crawler, and your `llms.txt` gives them a single, clean, factual document containing your services, your prices in OMR, your legal entity, your founder, your contact details and your URL map. It even pre-empts the brand collision with the unrelated `aiprofitlab.com` trading products. That file is doing real work, and it is better than what most GCC agencies have.

How LLMs actually decide who to cite

Three signals, roughly in this order of weight:

1. **Corroboration across independent sources.** A model trusts a fact it has seen in several places that are not you. This is the single biggest lever and the one you are weakest on.
2. **Extractability.** Can a specific, quotable claim be lifted cleanly from your page and attributed? Direct answers, definitions, numbers, and tables are extractable. Persuasive prose is not.
3. **Entity clarity.** Does the model know who you are, where you are, and what you do — unambiguously? This is where your brand collision hurts and where `llms.txt` and schema help.

THE UNCOMFORTABLE PART

Signal one cannot be produced on your own website. **No amount of on-site optimisation makes an LLM cite you if nothing off-site corroborates you.** Your two existing backlink documents are the right instinct, but treat their DR/traffic figures as unverified — they read as estimates rather than pulled data. The value in those documents is the target list, not the metrics.

The off-site programme — do these in this order

RANK	TARGET	WHY IT MATTERS HERE	WHAT TO ACTUALLY DO
1	Google Business Profile	Local pack + a Google-owned entity record	Verify it. Then post weekly, add services with prices, and answer your own FAQs in the Q&A section.
2	LinkedIn (founder + company)	Heavily crawled; strong entity anchor	Post the same insights as your articles, natively. Your operator background is the differentiator — use it.
3	Crunchbase	Feeds knowledge graphs directly	Create the entity, link Lotus Gulf International, CR 1570092, Muscat HQ.
4	Reddit + Quora	Disproportionately cited by AI search	Answer real Oman/GCC automation questions properly. No links in the first line, no pitch. Be the useful answer.
5	Omani business media	Local corroboration — the exact signal you lack	Pitch one Vision 2040 / SME-automation op-ed. Local beats global here.
6	GitHub org	High-trust, dofollow, genuinely on-brand	Publish something real and small: an Arabic/English WhatsApp prompt pack, or a PDPL compliance checklist.
7	Directories (Clutch, GoodFirms, Sortlist)	Entity consistency, some referral	Consistent NAP everywhere. Note most need client references you do not yet have — revisit after section 09.

On-site changes that raise extractability

- **Lead every article with a direct answer.** If the title asks a question, answer it in the first 40–60 words, plainly, before any context. That block is what gets lifted and quoted.
- **Publish one definition page per core term** — "AI receptionist", "WhatsApp Business API", "agentic automation", each with an Oman-specific framing. Definition pages are cited far out of proportion to their length.
- **Keep your prices on the page in text.** You already do this, and it is a real advantage: most competitors hide pricing, so when a model is asked what AI automation costs in Oman, you are one of the few extractable sources. Never move to "contact for pricing."
- **Add `knowsAbout` and widen `areaServed`** on the organisation schema — see Appendix A.
- **Date everything visibly.** You do this well. Keep `dateModified` honest — models weight freshness and will notice inconsistency.
- **Keep `llms.txt` synchronised with your prices.** It currently matches. The moment it drifts, you are feeding models stale figures that will be quoted back at you by prospects.

A MEASUREMENT YOU SHOULD RUN MONTHLY

Ask ChatGPT, Claude, Perplexity and Gemini the same five questions each month and log the answers verbatim: "Who does AI automation for small businesses in Oman?" · "How much does a WhatsApp AI receptionist cost in Oman?" · "Best AI automation company in Muscat" · "من يقدم أتمتة الذكاء الاصطناعي في عُمان؟" · "AI Profit Lab reviews". Track whether you are named, whether you are cited with a link, and — critically — whether you are being confused with the unrelated `aiprofitlab.com`. This is your GEO scoreboard and it takes fifteen minutes a month.

07 Ranking #1 in Google

Being honest about what "rank 1" means: for broad terms like "AI automation" it is not achievable and not worth chasing. For the specific, high-intent, Oman-anchored terms your buyers actually type, it is very achievable — and several are already close.

The three tiers

TIER	REALISTIC?	APPROACH
Brand "AI Profit Lab"	OWN IT NOW	You rank, but you share the name with unrelated companies. Entity disambiguation (section 06) is the whole job here.
Local commercial "AI automation company Muscat"	WINNABLE IN 3–6 MO	This is the money tier and where you currently lose. Needs commercial pages + verified GBP + proof.
Long-tail informational "WhatsApp AI receptionist cost Oman"	ALREADY WINNING	Defend and refresh. Do not write more of these until the commercial tier is built.
Broad head terms "AI automation"	NO	Global SaaS budgets own these. Ignore entirely.

STOP WRITING ARTICLES FOR A WHILE

This will feel counter-intuitive given the cadence you have kept, but: **you have enough articles**. 308 pieces at a 1,713-word median is a serious library, and adding a 155th English article has far lower marginal value than publishing your first real commercial page or your first case study. Redirect the writing time. Resume the cadence at roughly two per month in Q2, focused on refreshing and consolidating rather than expanding.

The commercial pages to build — in order

- "AI automation companies in Oman" — an honest category page that lists the real players, including competitors. Counter-intuitive, but category pages that name rivals earn links, get cited by LLMs, and rank because they genuinely serve the query. Draft copy already exists in `.agents/competitors/pages/01-...`.
- "AI automation vs hiring admin staff in Oman" — this is where your deals are actually won or lost, and your own research says so. The comparison is against a salary, not against QwicLink. Draft exists at `pages/02-...`.
- Industry landing pages — clinics, real estate, trading/distribution, car workshops. You already have articles on each; these are commercial pages with pricing, scope and a clear CTA, cross-linked from the matching article cluster.
- "WhatsApp AI receptionist Oman — pricing" — a dedicated commercial page. You rank for the informational version already; give the buying version its own page.
- Arabic equivalents of 1, 2 and 4 — lower competition, and you have committed to Arabic. Every one must go through a native Arabic reviewer before publishing. That is your own standing rule and it should not bend for speed.

HOLD THIS ONE BACK

The drafted **vs-QwicLink** page (`pages/04-...`) should **not** go live yet. A named-competitor page invites a side-by-side comparison, and in that comparison you are 2.3× more expensive with zero published clients against their published pricing. You would be sending buyers to a scoreboard you currently lose. Publish it once section 09 has produced two or three references.

Internal linking

A sampled article carries 19 unique internal links, of which only 4 point to money pages and 5 to other articles. That is thin for a 308-page library. Two changes:

- **Cluster the articles.** Group them into topic hubs — WhatsApp automation, PDPL/compliance, dashboards, industry guides — with each hub page linking down and each article linking back up. You have hub tooling already (`process_hubs.py` , `run_hubs.py`).
- **Point every article at the right commercial page**, not just at the generic services page. A clinic article should link to the clinic landing page.

08 Turning visitors into buyers

You asked how to turn every visitor into a buyer. You cannot — but you can stop losing the ones who were already interested, and that is where your losses currently are.

What happens today

A prospective buyer searches "how much does a WhatsApp AI receptionist cost in Oman", lands on your article, and reads a genuinely good 1,700-word piece. At the end they have three options: click a WhatsApp link and start a conversation with a stranger, navigate to a money page on their own initiative, or leave. Most leave — not because the content failed, but because the only offered action is the highest-commitment one on the page.

The gap in one line

Your funnel goes straight from "reading an article" to "message a human on WhatsApp". There is no middle step — no agent, no email capture, no low-commitment offer. Everyone not ready for WhatsApp today is lost.

The three-step fix

1. Aiden on every article — the zero-commitment way to ask a question.
2. One lead magnet — the low-commitment way to stay in touch.
3. A named next step at the end of every article — not a generic link.

Step one — Aiden everywhere

This is finding #1 and the highest-return change in the document. Beyond the direct conversion gain, it is the single best demonstration of your product you will ever run: a buyer reading about an AI agent that answers questions at 2am can, in that moment, ask one and watch it answer. Nothing on your services page persuades as well as that does.

Three details worth getting right:

- Initialise on idle or first scroll so it never delays the article's first paint.
- Make its opening line context-aware. On a pricing article: *"Want me to price this for your business?"* beats a generic greeting substantially.
- Make sure `aiden-index.json` covers the article content, so it can answer about what the visitor is reading. Your notes flag that prices do not auto-sync — verify before rollout, because an agent quoting a stale price on a pricing article is worse than no agent.

Step two — one lead magnet, offered well

You need exactly one, and it should be the thing you already reference: **the Silent Buyer Test**. It is concrete, it is diagnostic, it produces a personalised result, and it naturally leads to a build conversation. Offer it inline in the body of relevant articles, not as a popup, and not only in the footer.

What it should do: the visitor gives a business name and WhatsApp number; you run a short scripted test of how their business currently responds to an after-hours enquiry; they receive a one-page result. That result *is* the sales conversation, and it costs a prospect nothing to accept.

WHY THIS BEATS A PDF DOWNLOAD

A generic ebook collects addresses from people who will never buy. The Silent Buyer Test self-selects for owners who suspect they are losing enquiries — which is exactly your buyer — and it produces the evidence that makes your pitch for you. It also, conveniently, generates the raw material for section 09.

Step three — fix the friction on the money pages

PAGE	OBSERVED	RECOMMENDATION
/en/checkout/	16 form fields	Too many for a WhatsApp-first market. Cut to the minimum needed to start a conversation — name, business, WhatsApp, which system. Collect the rest in conversation.
/en/contact/	508 words, 0 CTA buttons	The thinnest money page on the site and it has no primary action styled as a button. Give it one obvious next step.
/ homepage	3 CTAs, no capture	CTAs are good and well-written. Add the Silent Buyer Test as a secondary path for visitors not ready to message.
/en/services/	13 CTAs	Well structured. Prices published — keep it that way. Add proof blocks as section 09 delivers them.

What not to do

- **No exit-intent popups.** They are disproportionately disliked in this market and would undercut the calm, operator-not-agency tone that is your actual differentiator.
- **Do not add a phone-first CTA.** WhatsApp-first is correct for Oman; your instincts here are right.
- **Do not fabricate urgency.** If the founding-price seat counter is real, show it. If it is not, remove it — a buyer who discovers invented scarcity discounts everything else you have said.

09 The Proof Engine

This is the load-bearing section. You have zero publishable proof, and that single fact caps the return on every other page in this document.

Zero of 375 pages carry a review, a rating, a named client or a verifiable outcome. Meanwhile a direct competitor publishes six named clients including Oman Flour Mills. When a buyer compares you side by side, the comparison is not about features — it is about who else has trusted you, and right now your answer is silence.

THE ONE RULE

Do not invent proof. Not a stock-photo testimonial, not a plausible-sounding case study, not `aggregateRating` schema with no reviews behind it. Fabricated review markup is a manual-action risk with Google, but the real cost is worse: in a market the size of Muscat's business community, being caught inventing a client is not a setback, it is the end. Everything below produces *real* proof.

Five sources of proof you can start this month

1. Proof from your own operation — available immediately

You do not need a client to prove the mechanism works. You need a running system, and you have one. Publish live, honest numbers from Aiden itself: how many questions it answered last month, median response time, the share handled without you, the split between Arabic and English. A public "Aiden this month" panel is proof of capability, it is fully truthful, it is unusual enough to earn links, and it demonstrates the exact product you sell. Start logging now so the first month's figures exist.

2. The Silent Buyer Test — proof as a by-product

Every test you run produces a real, anonymised data point: *"of 30 Muscat businesses tested, 22 never replied to an after-hours enquiry; median reply time was 14 hours."* That aggregate statistic is publishable without naming anyone, is genuinely newsworthy locally, is exactly the kind of claim LLMs cite, and it makes your core argument with evidence instead of assertion. Twenty to thirty tests is enough to publish.

3. Founding Partner pilots — with consent designed in

Your pilot mechanism already exists. What is missing is that consent and measurement are not built into it. From the next pilot onward, make three things part of the agreement itself:

- A **baseline measurement** taken before anything is built — response times, missed enquiries, hours spent. Without a before, there is no after.
- A **written consent clause** covering how they may be referenced, with tiered options: named with logo, named without figures, or anonymised by sector. Most people accept one of the three when asked properly at the start, and almost nobody agrees when asked cold six months later.
- A **90-day review** booked at kickoff, where you capture the numbers and ask for one or two sentences in their own words.

4. Demonstration in place of testimony

Until named clients exist, your demos carry the proof load. Your `/en/demos/` page exists and is schema-marked, but it is not doing enough work. Record short, unedited screen captures: a real Arabic enquiry arriving and being answered; the dashboard updating. Unedited and slightly rough beats polished — polish reads as marketing, roughness reads as real.

5. Third-party signals that are not testimonials

Google Business Profile reviews are the highest-value form here, because they are visible in the local pack and independently hosted. They can come from anyone with genuine first-hand experience — a pilot participant, someone you ran a Silent Buyer Test for, a business you advised. Ask directly and specifically. Five real reviews on a verified profile would materially change your local position.

The sequence

WHEN	PROOF ASSET	WHERE IT GOES
Month 1	Aiden operating statistics	Homepage panel, services page, demos page
Month 1–2	Silent Buyer Test aggregate (20–30 businesses)	Its own page + pitched to Omani business media
Month 2	Unedited demo recordings	Demos page, industry landing pages, LinkedIn
Month 2–3	First GBP reviews	Google local pack; quoted on site once genuine
Month 3+	First pilot case study with baseline vs. result	Case study page, Review schema, vs-competitor pages unlocked

WHAT THIS UNLOCKS

Once two or three references exist you can publish **Review** and **aggregateRating** schema truthfully, launch the named-competitor pages currently held back, complete the Clutch and GoodFirms listings that require client references, and answer the one question that is currently ending your deals.

10 Local SEO and Google Business Profile

Your Google Business Profile is the highest-leverage unclaimed asset you have, and it is currently blocked on one administrative step.

The address has been corrected to South Al Khuwair, Bousher — the CR address — after previously sitting in the Saudi desert roughly 60 km from Riyadh. But the profile is **not verified**, and until it is, it cannot appear in the Muscat local pack or on Maps. For a business whose entire positioning is "an operator in Muscat, not an offshore agency", that is the wrong thing to be missing.

Sequence

1. **Request verification this week.** It has a lead time measured in weeks, so it must start before anything else in the roadmap. Everything below is blocked on it.
2. **Reconcile the map pin with your schema.** Your schema currently uses `23.5803, 58.4310`, which is the OpenStreetMap centroid of the district rather than a surveyed office location. Once GBP is verified, take the pin coordinates from GBP and update the schema to match. **GBP wins** — the schema should follow it, not the other way round.
3. **Complete the profile fully.** Services with prices in OMR, service areas covering Muscat and the governorates you serve, hours, both Arabic and English descriptions, and real photos — of you, of the work, not stock imagery.
4. **Use the Q&A section deliberately.** You may ask and answer your own questions there. Seed the five questions buyers actually ask — cost, timeline, whether a monthly fee is required, whether it works in Arabic, what happens if they cancel. This content surfaces in the local pack and is read by AI search.
5. **Post weekly.** Short posts, linked to your articles. Low effort, and it keeps the profile active in a market where most competitors let theirs go stale.
6. **Build NAP consistency.** Every directory, profile and citation must carry exactly the same string: **Lotus Gulf International · CR 1570092 · South Al Khuwair, Bousher, Muscat · +968 9924 5250**. Inconsistent NAP is a common and entirely avoidable local-ranking drag — and your records show the site previously used a different entity name and district, so old citations with the outdated details may still be out there. Find and correct them.

WORTH KNOWING

Local pack results are one of the strongest inputs into "near me" and city-qualified AI answers. A verified, complete, actively-posted GBP with genuine reviews therefore serves both of your stated priorities at once — Google ranking and LLM citation. It is the highest ratio of result to effort in this entire document.

11 The 90-day roadmap

Thirteen weeks, sized to the 10–20 hours per week you have. Hour estimates are deliberately conservative. If a week runs short, the overflow task is always content — never proof, and never the technical blockers.

PHASE STRUCTURE

Weeks 1–2 — Unblock. Fix what is actively costing money and start the things with long lead times.

Weeks 3–6 — Connect. Wire the traffic to the selling machine.

Weeks 7–10 — Commercial. Build the pages that capture buying intent.

Weeks 11–13 — Compound. Off-site authority and the first real proof.

Phase 1 · Weeks 1–2 — Unblock

WEEK 1 · DO FIRST

~2 hrs

Start the Google Business Profile verification

Longest lead time in the plan, so it goes first. Request verification, then complete the profile while waiting: services with OMR prices, service areas, hours, EN + AR descriptions, real photos.

WEEK 1

~1 hr

301 www → apex

Appendix A. Then confirm with `curl -I https://www.aiprofitlab.io/` and expect a `301`, not the `200` it returns today.

WEEK 1–2 · HIGHEST RETURN

~6 hrs

Put Aiden on all 308 articles

Extend the widget tooling to the `blog/en/` and `blog/ar/` trees. Initialise on idle. Verify `aiden-index.json` carries current prices *before* rollout. Deploy, then confirm live on three random articles.

WEEK 2

~3 hrs

Set up the proof-capture habit

Start logging Aiden statistics. Write the consent clause and the baseline-measurement checklist into your pilot onboarding. This costs almost nothing now and is worth everything in month three.

WEEK 2

~2 hrs

Breadcrumb schema on money pages + image attributes

Appendix A. Small, fast, and improves how your commercial pages appear in results.

WEEK 2 · BASELINE

~2 hrs

Record your starting position

Run the five-question LLM check from section 06 and save the answers. Note current Search Console impressions and positions. You cannot show progress without a before.

Phase 2 · Weeks 3–6 — Connect

WEEK 3–4

~8 hrs

Build and launch the Silent Buyer Test

A short form, a scripted test procedure, and a one-page result template. Offer it inline on the twenty highest-traffic articles. This is your email capture, your lead qualifier and your proof generator in one asset.

WEEK 4

~4 hrs

Reduce checkout friction

Cut the 16 form fields to the four that matter. Give `/en/contact/` a proper primary CTA and more substance than its current 508 words.

WEEK 5

~6 hrs

Analytics replacement + money-page CSS

Swap GA4 for a lightweight analytics script — a 98% cut to your heaviest asset. Move money-page CSS to a shared hashed stylesheet using the pattern your blog already uses, keeping ~10 KB critical CSS inline.

WEEK 5-6

~8 hrs

Topic clusters and internal links

Group the 308 articles into hubs — WhatsApp automation, PDPL and compliance, dashboards, industry guides. Each article links up to its hub and across to the right commercial page. Use the existing hub tooling.

WEEK 6 · ONGOING FROM HERE

~4 hrs

Run the first Silent Buyer Tests

Target 20–30 Muscat businesses over the following weeks. Log every result — this becomes a publishable dataset.

Phase 3 · Weeks 7–10 — Commercial

WEEK 7-8

~10 hrs

Publish "AI automation companies in Oman"

The honest category page, naming real competitors. Draft copy exists in `.agents/competitors/pages/`. This earns links and citations precisely because it is not self-serving.

WEEK 8-9

~8 hrs

Publish "AI automation vs hiring admin staff"

The comparison your buyers are actually making. Argue against a salary, not against QwicLink. Draft exists.

WEEK 9-10

~12 hrs

Two industry landing pages

Pick the two sectors where you have the most article coverage and the clearest story — clinics and trading/distribution are the obvious candidates. Commercial pages with scope, pricing and CTA, linked from the matching article cluster.

WEEK 10

~6 hrs

Arabic commercial page — with native review

Translate and adapt the strongest of the above. It does not publish until a native Arabic speaker has reviewed it. Budget the review time realistically.

Phase 4 · Weeks 11–13 — Compound

WEEK 11

~6 hrs

Publish the Silent Buyer Test findings

Your first genuinely original, citable dataset. Give it its own page, then pitch it to Omani business media as a local story with real numbers.

WEEK 11–12

~8 hrs

Off-site authority push

Crunchbase entity. GitHub org with one genuinely useful public asset. Weekly LinkedIn posts. Substantive Reddit and Quora answers on Oman automation questions. Consistent NAP everywhere.

WEEK 12

~5 hrs

Publish the Aiden operating panel + demo recordings

Three months of real statistics now exist. Put them on the homepage and demos page, with unedited demo captures alongside.

WEEK 13 · REVIEW

~5 hrs

Re-measure and re-plan

Re-run the five LLM questions and compare to your week-2 baseline. Review Search Console. Confirm GBP is verified and gathering reviews. Decide whether the proof position now supports the vs-competitor pages.

IF YOU ONLY GET THROUGH HALF OF THIS

Do weeks 1–2 and weeks 3–4 completely. Aiden on the articles, the www redirect, GBP verification started, and the Silent Buyer Test live. Those four things alone address the largest conversion loss, the largest technical duplication, the local-search blocker, and the proof gap. Everything after that compounds — but nothing compounds if these four are not done.

12

Months 4–12, and how to measure it

After the sprint

- **Months 4–6.** First case study from a completed pilot, with baseline-versus-result numbers. This unlocks **Review** schema, the held-back competitor pages, and the directory listings requiring client references. Resume writing at roughly two articles per month, aimed at gaps rather than volume.
- **Months 6–9.** With proof in hand, complete Clutch, GoodFirms and Sortlist. Pitch one substantial op-ed to regional business media. Consider the URL restructure (finding #10) if it still matters to you — this is the safe window for it, not the sprint.
- **Months 9–12.** Expand the industry landing pages to the remaining sectors. Build the Arabic commercial set out fully. Reassess whether GCC expansion beyond Oman is worth the dilution — on current evidence, deepening in Oman beats widening.

What to measure — and what to ignore

METRIC	CADENCE	WHY IT EARNS ITS PLACE
WhatsApp conversations started	Weekly	The closest measurable thing to revenue. This is the number that matters most.
Aiden conversations + questions asked	Weekly	Proves the article rollout worked, and doubles as publishable proof.
Silent Buyer Tests requested / run	Weekly	Your lead pipeline and your proof pipeline in one number.
LLM citation check (5 questions)	Monthly	Your GEO scoreboard. Fifteen minutes, and the only honest read on your stated top priority.
GBP: views, calls, direction requests	Monthly	The clearest signal of local intent reaching you.
Search Console: commercial-query impressions	Monthly	Filter to buying-intent terms. Total impressions are flattered by informational traffic you already win.
Referring domains	Monthly	The off-site corroboration that drives LLM citation.
Total pageviews	—	Ignore. It rises with article count and tells you nothing about whether anyone is buying.
Lighthouse score	—	Check once after the week-5 work, then stop. You are already fast enough.

THE REALISTIC OUTCOME

If the 90 days are executed as written: expect materially more conversations from the traffic you already have, a verified local presence you currently lack, commercial pages competing for buying queries for the first time, and — most importantly — the first real proof assets in the company's history. What you should **not** expect is a step-change in traffic volume. You do not have a traffic problem. You have a conversion problem, a proof problem and a commercial-page problem, and this plan is aimed squarely at those three.

A Appendix A — copy-paste fixes

A1 · 301 redirect www → apex

Add near the top of `public_html/.htaccess`, before other rewrite rules. Your file already has a well-organised rewrite section — place it there.

```
# --- Canonical host: www -> apex (301) ---
# www.aiprofitlab.io currently serves HTTP 200, duplicating all 332 URLs.
# Canonicals already point at the apex; this makes it authoritative.
<IfModule mod_rewrite.c>
  RewriteEngine On
  RewriteCond %{HTTP_HOST} ^www\.aiprofitlab\.io$ [NC]
  RewriteRule ^(.*)$ https://aiprofitlab.io/$1 [R=301,L]
</IfModule>
```

Verify after deploying:

```
curl -sI https://www.aiprofitlab.io/ | head -1
# expect: HTTP/2 301 (today it returns HTTP/2 200)
```

IF THE REDIRECT DOES NOT TAKE EFFECT

Hostinger's CDN may already hold the `200` at the edge. Purge the CDN cache after deploying, then re-test. A redirect that works at origin but not at the edge is a known failure mode on this stack.

A2 · BreadcrumbList for money pages

Add to the existing `@graph` on each money page. Example for `/en/services/`:

```
{
  "@type": "BreadcrumbList",
  "@id": "https://aiprofitlab.io/en/services/#breadcrumb",
  "itemListElement": [
    { "@type": "ListItem", "position": 1,
      "name": "Home", "item": "https://aiprofitlab.io/" },
    { "@type": "ListItem", "position": 2,
      "name": "What I build", "item": "https://aiprofitlab.io/en/services/" }
  ]
}
```

Apply the same shape to `/en/process/`, `/en/about/`, `/en/contact/`, `/en/checkout/` and `/en/demos/`, and to the Arabic equivalents at `/services/`, `/process/`, `/about/` and `/contact/` with Arabic `name` values. Generate it in `tools/build_v4.py` rather than hand-editing, so it survives the next rebuild.

A3 · Widen the organisation entity

Two additions to the `ProfessionalService` node. `areaServed` currently lists Oman only; `knowsAbout` is absent entirely and is one of the strongest topical-authority signals available.

```

"areaServed": [
  { "@type": "Country", "name": "Oman" },
  { "@type": "Country", "name": "United Arab Emirates" },
  { "@type": "Country", "name": "Saudi Arabia" },
  { "@type": "Country", "name": "Qatar" },
  { "@type": "Country", "name": "Kuwait" }
],
"knowsAbout": [
  "AI automation for small and medium businesses",
  "WhatsApp Business API integration",
  "Arabic and English conversational AI agents",
  "Business process automation",
  "Oman Personal Data Protection Law compliance",
  "Sales and operations dashboards",
  "Quote and invoice follow-up automation"
]

```

ONE CAUTION

Only widen `areaServed` to countries you will genuinely take work in. It is an entity claim, and claiming markets you do not serve muddies exactly the entity clarity you are trying to build. If Oman is truly the focus, list Oman and the one or two neighbours you would actually accept a project from.

A4 • Image attributes

For the five homepage images below the fold. Do **not** apply this to the two preloaded hero posters — those are correctly configured and lazy-loading them would slow your largest contentful paint.

```



```

A5 • Defer Google Analytics (if you keep it)

Preferred path is replacement with a lightweight script. If GA4 stays, load it out of the critical path:

```

// Loads gtag only when the browser is idle, or on first interaction.
// Removes ~188 KB from the critical path.
(function () {
  var loaded = false;
  function loadGA() {
    if (loaded) return; loaded = true;
    var s = document.createElement('script');
    s.async = true;
    s.src = 'https://www.googletagmanager.com/gtag/js?id=G-SLR9GD3MJP';
    document.head.appendChild(s);
    window.dataLayer = window.dataLayer || [];
    function gtag(){ dataLayer.push(arguments); }
    gtag('js', new Date());
    gtag('config', 'G-SLR9GD3MJP');
  }
  if ('requestIdleCallback' in window) requestIdleCallback(loadGA, { timeout: 5000 });
  else setTimeout(loadGA, 3000);
  ['scroll','pointerdown','keydown'].forEach(function (e) {
    window.addEventListener(e, loadGA, { once: true, passive: true });
  });
})();

```

A6 • Verification commands

```
# Aiden present on articles? (expect 1, currently 0)
curl -s https://aiprofitlab.io/blog/en/2026-07-10-what-is-an-ai-agent-oman/ \
  | grep -c aiden-chat.js

# www canonical host (expect 301)
curl -sI https://www.aiprofitlab.io/ | head -1

# Breadcrumbs on a money page (expect 1)
curl -s https://aiprofitlab.io/en/services/ | grep -c BreadcrumbList

# Confirm llms.txt prices still match the services page
curl -s https://aiprofitlab.io/llms.txt | grep OMR
```

PACE YOUR CHECKS

Space these a few seconds apart. Rapid bursts trip Hostinger's rate limiting and return **403**, which looks alarming but is not a site fault — as this audit discovered first-hand.

B Appendix B — keyword targets

ON SEARCH VOLUMES

No volume figures are given here, because I do not have verified keyword data for Oman and estimated numbers for a market this size are unreliable enough to mislead. Terms are ranked by **buying intent and winnability** instead, which is the more useful ordering for you. If you want volumes, pull them from Search Console — you already rank for many of these, so your own impression data is better than any third-party estimate.

Tier 1 · Commercial — build pages for these

QUERY	STATUS	DESTINATION
ai automation company oman	COMPETITORS OWN	Category page (roadmap wk 7–8)
ai automation muscat	COMPETITORS OWN	Category page + GBP
whatsapp automation oman	COMPETITORS OWN	WhatsApp commercial page
whatsapp ai receptionist oman price	ARTICLE RANKS	Dedicated pricing page
ai chatbot arabic english oman	PARTIAL	WhatsApp commercial page
business automation consultant muscat	GAP	Services page + GBP
ai for clinics / real estate / trading oman	ARTICLES ONLY	Industry landing pages (wk 9–10)

Tier 2 · Comparison — high intent, needs proof first

QUERY	STATUS	DESTINATION
ai automation vs hiring staff oman	BUILD NOW	Draft exists — publish wk 8–9
best ai automation companies in oman	BUILD NOW	Honest category page
qwiclink alternative / vs	HOLD	Blocked until proof exists (§09)
nuqta alternative / vs	HOLD	Blocked until proof exists (§09)

Tier 3 · Informational — already winning, defend only

You rank for these and should not write more of them this quarter. Refresh dates, tighten the opening answer for extractability, and make sure each links to the right commercial page.

QUERY	NOTE
what is an ai agent oman	Ranking. Add a 40-word direct answer at the top.
whatsapp ai receptionist cost oman	Ranking. Link to the new pricing page when it exists.
oman pdpl ai chatbot compliance	Strong cluster. Genuine differentiator — few competitors cover it.
whatsapp business api vs app oman	Ranking. Good top-of-funnel entry point.
ai automation oman vision 2040	Local-authority topic. Useful for media pitches.

Tier 4 • Arabic — lower competition, needs native review

QUERY	NOTE
أتمتة الذكاء الاصطناعي عمان	Core commercial term. Arabic category page.
شات بوت واتساب عمان	WhatsApp commercial intent.
موظف استقبال ذكي واتساب	Maps to your strongest English cluster.
تكلفة أتمتة الأعمال في عمان	Pricing intent — and you publish prices, which most rivals do not.
شركات الذكاء الاصطناعي في مسقط	Local commercial. Pairs with GBP.

STANDING RULE

No Arabic page ships without review by a native Arabic speaker. This is your own rule and it exists for good reason — machine-drafted Arabic marketing copy reads as foreign to exactly the buyers you are trying to reach, and in this market that costs more trust than the ranking is worth.

Nahid Abyari • Founder, AI Profit Lab
hello@aiprofitlab.io • +968 9924 5250 • aiprofitlab.io
AI Profit Lab is a brand of Lotus Gulf International • CR 1570092 • South Al Khuwair, Bousher, Muscat, Oman
All findings verified against the live site on 25 August 2026. Competitor pricing and third-party authority metrics are as recorded in internal research and should be re-checked before being quoted externally.